

# **TTIC 31230, Fundamentals of Deep Learning**

David McAllester, Autumn 2020

## **Actor-Critic Algorithms and AC3**

# REINFORCE

$$\nabla_{\Phi} E_{\pi_{\Phi}} R = E_{\pi_{\Phi}} \sum_{t, t' \geq t} (\nabla_{\Phi} \ln \pi_{\Phi}(a_t | s_t)) r_{t'}$$

Sampling runs and computing the above sum over  $t$  and  $t'$  is Williams' REINFORCE algorithm.

## The Variance Issue

REINFORCE typically suffers from high variance of the gradient samples requiring very small learning rates and very long convergence times.

$$\nabla_{\Phi} E_{\pi_{\Phi}} R = \sum_{t, t' \geq t} E_{s_t, a_t, r_{t'}} (\nabla_{\Phi} \ln \pi_{\Phi}(a_t | s_t)) r_{t'}$$

We have to consider

- the variation over  $r_{t'}$  given  $s_t, a_t$
- the variation over the choice of  $s_t, a_t$

## Reducing the Variance over $r_{t'}$

### The Policy Gradient Theorem

“Policy Gradient Methods for Reinforcement Learning with Function Approximation” Sutton, McAllester, Singh, Mansour, 2000, cited by 2,841 as of March 2020.

“Actor-Critic Algorithms”, Konda and Tsitsilas, 2000, cited by 776.

These two papers both appeared at NeurIPS 2000 and are essentially identical. The first is just easier to read.

## Reducing the Variance over $r_{t'}$

$$\begin{aligned}
 \nabla_{\Phi} E_{\pi_{\Phi}} R &= \sum_{t,t'} E_{s_t,a_t,r_{t'}} \nabla_{\Phi} \ln \pi_{\Phi}(a_t|s_t) \textcolor{red}{r_{t'}} \\
 &= \sum_t E_{s_t,a_t} \nabla_{\Phi} \ln \pi_{\Phi}(a_t|s_t) \textcolor{red}{\sum_{t' \geq t} E_{r_{t'} | s_t,a_t} r_{t'}} \\
 &= E_{\pi_{\Phi}} \sum_t (\nabla_{\Phi} \ln \pi_{\Phi}(a_t|s_t)) \textcolor{red}{Q^{\pi_{\Phi}}(s_t, a_t)}
 \end{aligned}$$

$$Q^{\pi}(s, a) = E_{\pi} \sum_t r_t \mid s_0 = s, a_0 = a$$

## Reducing the Variance over $r_{t'}$

$$\nabla_{\Phi} E_{\pi_{\Phi}} R = \sum_t E_{s_t, a_t} (\nabla_{\Phi} \ln \pi_{\Phi}(a_t|s_t)) \textcolor{red}{Q}^{\pi_{\Phi}}(\textcolor{red}{s}_t, \textcolor{red}{a}_t)$$

**The point** is that we can now approximate  $Q^{\pi_{\Phi}}$  with neural network  $Q_{\Phi}$  where the networks  $\pi_{\Phi}$  and  $Q_{\Phi}$  can use different, perhaps overlapping, parts of  $\Phi$ .

We reduced the variance at the cost of approximating the expected future reward.

# The Actor-Critic Algorithm

$$\nabla_{\Phi} E_{\pi_{\Phi}} R \approx E_{\pi_{\Phi}} \sum_t (\nabla_{\Phi} \ln \pi_{\Phi}(a_t|s_t)) Q_{\Phi}(s_t, a_t)$$

$\pi_{\Phi}$  is the “actor” and  $Q_{\Phi}$  is the “critic”

## The Actor-Critic Algorithm

$$\nabla_{\Phi} E_{\pi_{\Phi}} R \approx E_{\pi_{\Phi}} \sum_t (\nabla_{\Phi} \ln \pi_{\Phi}(a_t|s_t)) Q_{\Phi}(s_t, a_t)$$

We can sample an episode and then do

$$\begin{aligned} \Phi & += \sum_t \eta_1 (\nabla_{\Phi} \ln \pi_{\Phi}(a_i|s_i)) Q_{\Phi}(s_t, a_t) \\ \Phi & -= \sum_t \eta_2 \nabla_{\Phi} \left( Q_{\Phi}(s_t, a_t) - \sum_{t' \geq t} r_{t'} \right)^2 \end{aligned}$$

The two updates typically apply to different (but perhaps overlapping) subsets of the parameters  $\Phi$ .



## Reducing the Variance over $s_t$

Thorem:

$$\nabla_{\Phi} E_{\pi_{\Phi}} R = \sum_t E_{s_t, a_t} (\nabla_{\Phi} \ln \pi_{\Phi}(a_t|s_t)) (Q^{\pi_{\Phi}}(s_t, a_t) - V^{\pi_{\Phi}}(s_t))$$

$$V^{\pi_{\Phi}}(s) = E_{a \sim \pi_{\Phi}(a|s)} Q^{\pi_{\Phi}}(s, a)$$

$Q^{\pi_{\Phi}}(s, a) - V^{\pi_{\Phi}}(s)$  is the “advantage” of deterministically using  $a$  rather than sampling an action.

## Proof

We have the following for any function  $V(s)$  of states.

$$\begin{aligned} & E_{s_t, a_t} (\nabla_{\Phi} \ln \pi_{\Phi}(a_t|s_t)) V(s_t) \\ &= E_{s_t} \sum_{a_t} (\pi_{\Phi}(a_t|s_t) \nabla_{\Phi} \ln \pi_{\Phi}(a_t|s_t)) V(s_t) \\ &= E_{s_t} \sum_{a_t} (\nabla_{\Phi} \pi_{\Phi}(a_t|s_t)) V(s_t) \\ &= E_{s_t} V(s_t) \nabla_{\Phi} \sum_{a_t} \pi_{\Phi}(a_t|s_t) = 0 \end{aligned}$$

## Advantage-Actor-Critic Algorithm

$$\nabla_{\Phi} E_{\pi_{\Phi}} R \approx E_{\pi_{\Phi}} \sum_t (\nabla_{\Phi} \ln \pi_{\Phi}(a_t|s_t)) (Q_{\Phi}(s_t, a_t) - V_{\Phi}(s_t))$$

We can sample an episode and then do

$$\begin{aligned}\Phi & += \sum_t \eta_1 (\nabla_{\Phi} \ln \pi_{\Phi}(a_i|s_i)) (Q_{\Phi}(s_t, a_t) - V_{\Phi}(s_t)) \\ \Phi & -= \sum_t \eta_2 \nabla_{\Phi} \left( Q_{\Phi}(s_t, a_t) - \sum_{t' \geq t} r_{t'} \right)^2 \\ \Phi & -= \sum_t \eta_3 \nabla_{\Phi} (V_{\Phi}(s_t) - Q_{\Phi}(s_t, a))^2\end{aligned}$$

## Asynchronous Methods for Deep RL (A3C)

Mnih et al., Arxiv, 2016 (Deep Mind)

$\tilde{\Phi} = \Phi$  (retrieve global  $\Phi$ )

using policy  $\pi_{\tilde{\Phi}}$  compute  $s_t, a_t, r_t, \dots, s_{t+K}, a_{t+K}, r_{t+K}$

$$R_i = \sum_{\delta=0}^D \gamma^{i+\delta} r_{(i+\delta)}$$

$$\Phi \ += \ \eta \sum_{i=t}^{t+K-D} \left( \nabla_{\tilde{\Phi}} \ln \pi_{\tilde{\Phi}}(a_i | s_i) \right) \left( R_i - V_{\tilde{\Phi}}(s_i) \right)$$

$$\Phi \ -= \ \eta \sum_{i=t}^{t+K-D} \nabla_{\tilde{\Phi}} \left( V_{\tilde{\Phi}}(s_i) - R_i \right)^2$$

## Issue: Policies must be Exploratory

The optimal policy is deterministic —  $a(s) = \operatorname{argmax}_a Q(s, a)$ .

However, a deterministic policy never samples alternative actions.

Typically one forces a random action some small fraction of the time.

## **Issue: Discounted Reward**

DQN and A3C use discounted reward on episodic or long term problems.

Presumably this is because actions have near term consequences.

This should be properly handled in the mathematics, perhaps in terms of the mixing time of the Markov process defined by the policy.

## Observation: Continuous Actions are Differentiable

In problems like controlling an inverted pendulum, or robot control generally, a continuous loss can be defined and the gradient of loss of with respect to a deterministic policy exists.

## More Videos

<https://www.youtube.com/watch?v=g59nSURxYgk>

<https://www.youtube.com/watch?v=rAai4QzcYbs>



**END**