

TTIC 31230, Fundamentals of Deep Learning

David McAllester, Autumn 2020

Deep Reinforcement Learning

Q -Learning

The Q Function

For discounted reward:

$$Q^\pi(s, a) = E_\pi \sum_t \gamma^t r_t \mid \pi, s_0 = s, a_0 = a$$

$$Q^*(s, a) = \sup_{\pi} Q^\pi(s, a)$$

$$\pi^*(a|s) = \operatorname{argmax}_a Q^*(s, a)$$

$$Q^*(s, a) = R(s, a) + \gamma E_{s' \sim P_T(\cdot|s, a)} \max_{a'} Q^*(s', a')$$

Q Function Iteration

It is possible to define Q -iteration by analogy with value iteration, but this is generally not discussed.

Value iteration is typically done for finite state spaces. Let S be the number of states and A be the number of actions.

One update of a Q table takes $O(S^2A^2)$ time while one update of value iteration is $O(S^2A)$.

Q -Learning

When learning by updating the Q function we typically assume a parameterized Q function $Q_\Phi(s, a)$.

Bellman Error:

$$\text{Bell}_\Phi(s, a) \doteq \left(Q_\Phi(s, a) - \left(R(s, a) + \gamma \mathbb{E}_{s' \sim P_T(s'|s, a)} \max_{a'} Q_\Phi(s', a') \right) \right)^2$$

Theorem: If $\text{Bell}_\Phi(s, a) = 0$ for all (s, a) then the induced policy is optimal.

Algorithm: Generate pairs (s, a) from the policy $\arg\max_a Q_\Phi(s_t, a)$ and repeat

$$\Phi \leftarrow \Phi - \eta \nabla_\Phi \text{Bell}_\Phi(s, a)$$

Issues with Q -Learning

Problem 1: Nearby states in the same run are highly correlated. This increases the variance of the cumulative gradient updates.

Problem 2: SGD on Bellman error tends to be unstable. Failure of Q_Φ to model unused actions leads to policy change (exploration). But this causes Q_Φ to stop modeling the previous actions which causes the policy to change back ...

To address these problems we can use a **replay buffer**.

Using a Replay Buffer

We use a replay buffer of tuples (s_t, a_t, r_t, s_{t+1}) .

Repeat:

1. Run the policy $\operatorname{argmax}_a Q_\Phi(s, a)$ to add tuples to the replay buffer. Remove oldest tuples to maintain a maximum buffer size.
2. $\Psi = \Phi$
3. for N times select a random element of the replay buffer and do

$$\Phi \leftarrow \Phi - \eta \nabla_\Phi (Q_\Phi(s_t, a_t) - (r_t + \gamma \max_a Q_\Psi(s_{t+1}, a)))^2$$

Replay is Off-Policy

Note that the replay buffer is from a **mixture of policies** and is **off-policy** for $\operatorname{argmax}_a Q_{\Phi}(s, a)$. This seems to be important for stability.

This seems related to the issue of stochastic vs. deterministic policies. More on this later.


Multi-Step Q -learning

$$\Phi \leftarrow \sum_t \nabla_{\Phi} \left(Q_{\Phi}(s_t, a_t) - \sum_{\delta=0}^D \gamma^{\delta} r_{(t+\delta)} \right)^2$$

Human-level control through deep RL (DQN)

Mnih et al., Nature, 2015. (Deep Mind)

We consider a CNN $Q_{\Phi}(s, a)$.



Watch The Video

<https://www.youtube.com/watch?v=V1eYniJ0Rnk>

Asynchronous Q-Learning (Simplified)

No replay buffer. Many asynchronous threads each repeating:

$\tilde{\Phi} = \Phi$ (retrieve Φ)

using policy $\operatorname{argmax}_a Q_{\tilde{\Phi}}(s, a)$ compute

$$s_t, a_t, r_t, \dots, s_{t+K}, a_{t+K}, r_{t+K}$$

$$\Phi \mathrel{:=} \eta \sum_{i=t}^{t+K-D} \nabla_{\tilde{\Phi}} \left(Q_{\tilde{\Phi}}(s_i, a_i) - \sum_{\delta=0}^D \gamma^{\delta} r_{i+\delta} \right)^2 \text{ (update } \Phi)$$

END